

# ALERT: Supersession-Aware Retrieval for Product-Litigation Risk Entailment

Anonymous ACL submission

## Abstract

Retrieval-augmented generation (RAG) treats semantically similar documents as interchangeable evidence, but in law a lexically close precedent may already have been narrowed, superseded, or overruled at the date that matters. We study this failure in a product-facing setting and introduce *Product-Litigation Risk Entailment* (PLRE), a directional task that asks whether an AI product’s attributes entail exposure to a legal-risk theory under precedent that is valid at a given query date. Our system, ALERT (Agentic Legal Exposure and Risk Tracking), infers *when* each authority stops being operative: it extracts treatment edges (overrules, supersedes, narrows) from later opinions, turns them into validity intervals, and validates them against gold annotations and a commercial citator (negative-treatment micro-F1 0.81). A deliberately simple retrieval-time gate applies these intervals before generation, and a goal-conditioned agentic loop stops when evidence coverage, not legal correctness, is satisfied. On ALERT-Dataset (1,247 U.S. Alitigation cases, 312 product-mapping cases), validity gating raises PLRE F1 from 0.72 to 0.78 and cuts stale citations from 12.7% to 4.1% relative to flat similarity, and ALERT improves over rule-based, single-pass LLM, RAG-only, single-agent, and iteration-matched self-refinement baselines on a human-adjudicated subset (macro F1 0.79 vs. 0.75 for the strongest baseline). We release the dataset with per-record label provenance, annotation guidelines, and evaluation scripts. ALERT is decision support for qualified reviewers, not legal advice.

## 1 Introduction

Generative-AI products increasingly face litigation over training data, output reproduction, model behavior, and downstream misuse. For the engineering organizations that build them, the operational question is not merely whether a lawsuit exists, but whether a theory of liability maps to a specific

product attribute, whether the supporting authority was valid at the relevant date, and which mitigation should be reviewed by counsel. Keyword alerts surface noisy filings, docket trackers rarely map cases to product features, and conventional retrieval-augmented generation (RAG) (Lewis et al., 2020) treats semantically similar legal documents as comparable even when later decisions have narrowed or superseded them.

We formalize this need as *Product-Litigation Risk Entailment* (PLRE): deciding whether product attributes entail exposure to a legal-risk theory under time-valid precedent. PLRE differs from standard legal textual entailment (Rabelo et al., 2022) in three ways: the premise is an engineering description rather than legal text, the relation is directional (product attribute  $\rightarrow$  risk theory), and the support set is conditioned on which authority was operative at the query date. Legal citators such as Shepard’s and KeyCite already encode the intuition that authority changes over time; we instead build a public-corpus, automatically extracted, time-indexed treatment layer and connect it to product-facing entailment rather than retrospective case lookup.

We present ALERT (**Agentic Legal Exposure and Risk Tracking**), a continuous sensing system built around this task. Its core component is a supersession-aware treatment layer: it infers when each authority stops being operative from treatment edges in *later* opinions, and a simple gate (Eq. (3)) applies this before generation, rather than flagging stale citations after an LLM has already produced a report. This matters because stale but lexically similar precedent can dominate ordinary RAG. Around this layer, a goal-conditioned agentic loop exposes unmet evidence requirements and stops when calibrated evidence coverage is satisfied. The layer’s form is not specific to law—it is defined for any time-stamped corpus in which later documents can supersede earlier ones—but we evaluate it only for law, within ALERT.

Our contributions are:

- a supersession-aware treatment layer: treatment edges extracted from public opinions, validated against gold spans and a commercial citator, and propagated into validity intervals that gate retrieval at query time (§4.4);
- PLRE, an initial product-conditioned legal-risk entailment benchmark with non-ALERT baselines and ablations that separate temporal validity from flat similarity (§4.5, §6);
- a goal-conditioned agentic loop whose stopping rule controls evidence coverage, with a conformal omission-risk guarantee (§3.1);
- ALERT-Dataset, with per-record label provenance, annotation guidelines, and a human-adjudicated evaluation subset (§5), plus a 30-day live deployment (Appendix D).

## 2 Background and Scope

**Legal-risk monitoring for AI products.** Copyright litigation against generative-AI systems has grown sharply since 2022, challenging web-scale training data (Henderson et al., 2023; Levendowski, 2018) under theories spanning DMCA §1202 information removal, direct infringement, and vicarious liability for outputs (Grimmelmann, 2016). Legal-operations teams track dockets, news, and internal product changes, but the bridge from a litigation event to a product-specific engineering action is usually manual. ALERT targets this product-facing layer. Its decision-support boundary is set out in the Ethics Statement.

## 3 Problem Formulation

Let  $\mathcal{D} = \{d_1, \dots, d_n\}$  be public litigation documents with metadata and text, and let  $P$  be a product specification describing architecture, data provenance, modality, and deployment context. ALERT returns a structured risk report  $R^*$  and mitigation candidates  $G^*$  that must satisfy an auditable goal predicate  $\mathcal{G}$  (§3.1). Report quality is not optimized directly; we evaluate it by severity-label F1 (RQ1), blinded report ratings (RQ2), and the coverage loss  $L_\tau$  (RQ6). The core difficulty is a domain gap: product specifications use engineering terms (data pipelines, API distribution), whereas litigation documents use causes of action, procedural posture, and precedent. PLRE (§4.5) therefore requires mapping engineering-language attributes to

legal-risk theories while conditioning the support set on authority valid at the query date.

### 3.1 Goal Predicate and Stopping Rule

ALERT treats report generation as goal-conditioned planning. The predicate is:

$$\mathcal{G}(R) = \mathcal{G}_{\text{risk}}(R) \wedge \mathcal{G}_{\text{evidence}}(R) \wedge \mathcal{G}_{\text{guide}}(R), \quad (1)$$

where  $\mathcal{G}_{\text{risk}}$  requires every risk item to have a severity label,  $\mathcal{G}_{\text{evidence}}$  requires at least  $k$  supporting documents above relevance threshold  $\tau$ , and  $\mathcal{G}_{\text{guide}}$  requires mitigation text for each High or Medium item. The controller chooses among SEARCH, ANALYZE, SUMMARIZE, MAP, VERIFY, and REPORT; the deterministic (LLM-free) Goal Checker either terminates, emits a structured critique to close residual gaps, or escalates unmet conditions. For example, if a High-risk item has only one valid authority after gating, the critique records the missing theory, product attribute, jurisdiction, query date, and required evidence count; the Planner rewrites this into narrower precedent and entity queries rather than simply retrying the same prompt.

Let  $\Phi(R_t)$  count the unmet atomic subconditions of  $\mathcal{G}$ , let  $\mathcal{T}_c^H$  be the human-labeled High-risk theories of case  $c$  that are *retrievable*, i.e., supported by at least one authority in  $\mathcal{D}$  that is valid at  $t_q$ , and define the coverage loss at evidence threshold  $\tau$  as

$$L_\tau(c) = \mathbf{1}[\exists T \in \mathcal{T}_c^H : T \notin R^*]. \quad (2)$$

**Proposition 1** (Coverage-controlled stopping with escalation). *Under monotone evidence caching and attainable progress,  $\Phi(R_t)$  is non-increasing, reaches zero in at most  $\Phi(R_0)$  successful progress steps when all atoms are attainable, and otherwise returns the residual atoms at  $T_{\max}$ . If  $\hat{\tau}$  is selected on exchangeable calibration cases by conformal risk control, then the next-case omission risk satisfies  $\mathbb{E}[L_{\hat{\tau}}] \leq \alpha$ .*

The first clause follows almost directly from the definitions: monotone caching means no satisfied atom becomes unmet, and attainable progress (each pass closes at least one attainable atom) is a design assumption that we do not verify directly. The second clause is a standard application of conformal risk control (Angelopoulos et al., 2023) to  $L_\tau$ , which is monotone in  $\tau$  (Appendix H). It covers evidence coverage only, holds only for exchangeable cases—our temporal split violates it at

$\alpha \in \{0.05, 0.10\}$  (§6)—and its tightness depends on the calibration-set size, which we do not report.

## 3.2 Risk Scoring

In the agentic layer, the Analyzer assigns each risk item a severity (High/Medium/Low),  $\text{RiskScore}(d_i, P, C)$ , by prompting the LLM backbone (App. G) with litigation document  $d_i$ , product specification  $P$ , and the precedent context  $C$  retrieved by LegalRisk-RAG (§4.2). The ingestion layer uses an additive scoring matrix (Table A.3, appendix) for real-time triage, with a critical-composite bonus when copyright and model-training evidence co-occur; this heuristic is treated as a baseline and never as ground truth.

# 4 The ALERT Framework

## 4.1 System Overview

ALERT has two layers. A sensing pipeline discovers AI-litigation filings in CourtListener/RECAP (Free Law Project, 2023) and legal news and turns them into structured case records; it is engineering infrastructure and is documented in Appendix D (architecture in Figure A.1). Below we describe the evaluated NLP components (§4.2–4.5).

## 4.2 LegalRisk-RAG

The LegalRisk-RAG pipeline adapts standard RAG to legal text in three ways. **Chunking** segments documents at *logical structural boundaries*—numbered claims, “WHEREFORE” clauses, section headings—rather than fixed token windows, and each chunk retains a metadata header (case id, filing date, court level, section type). **Embedding**: we compared a dense retriever (Karpukhin et al., 2020) (text-embedding-3-large), RECAP-tuned LegalBERT (Chalkidis et al., 2020), and a hybrid BM25 + neural re-ranker on a 200-pair development set; the hybrid (best) is the default (Table A.9, appendix). **Dynamic context expansion (precedent chaining)**: when a retrieved chunk references a prior case (e.g., *Authors Guild v. Google*), the pipeline appends that opinion as secondary context and queries the temporal precedent graph for subsequent treatment, so the Analyzer receives both semantic neighbors and validity metadata (cited positively, narrowed, superseded, or overruled at the filing date).

## 4.3 Goal-Conditioned Agentic Loop

The four specialized agents and the central Goal Checker operate as in Algorithm 2. The **Planner** decomposes unmet sub-goals into ordered retrieval/analysis sub-tasks; the **Retriever** runs LegalRisk-RAG queries and reformulates via legal entity term expansion when  $\mathcal{G}_{\text{evidence}}$  is unmet; the **Analyzer** computes  $\text{RiskScore}(d_i, P, C)$ , clusters evidence, and maps clusters to a risk taxonomy (Training Data, Output Reproduction, API Misuse, Model Weights); the **Goal Checker** evaluates  $\mathcal{G}(R_t)$  after each pass and returns machine-readable residuals such as `missing(valid-authority, DMCA, output-CMI, NDCA, 2025-02)`; the Planner converts these critiques into expanded entity/date/theory queries; and the **Reporter** assembles the risk register, per-risk evidence with citations, and mitigation guidance.

## 4.4 Supersession-Aware Temporal Precedent Graph

ALERT’s central component is a treatment layer that determines *when* each authority stops being operative; a simple retrieval-time gate then applies it. The layer is a temporal precedent graph  $\mathcal{P}_t = (V, E_t)$ . Each node  $v \in V$  denotes a case, order, or doctrinal rule and stores jurisdiction, court level, issue tags, filing/decision dates, and a validity interval  $[s_v, e_v)$ , where  $e_v$  is open unless the node is later vacated, overruled, superseded, or materially narrowed. Directed edges encode treatment relations: CITES, DISTINGUISHES, NARROWS, SUPERSEDES, and OVERRULES. Retrieval for a query at date  $t_q$  first obtains semantically relevant candidates and then applies validity gating:

$$\text{score}(v, q, t_q) = \text{sim}(v, q) \cdot \mathbf{1}[s_v \leq t_q < e_v] \cdot w(\text{treat}(v, t_q)). \quad (3)$$

Here  $w(\cdot)$  is a fixed, bounded monotone schedule in  $(0, 1]$ : positively treated or still-open precedents take weight 1, cautionary or materially narrowed authorities are down-weighted, and fully closed authorities are already removed by the indicator (the exact schedule is released with the implementation). This separates legal authority from topical similarity: a stale but lexically close case may still be shown as context, but it cannot support a risk decision without passing the validity check. Eq. (3) is thus a constrained retrieval operator over  $(q, t_q)$ : a host retriever scores semantic relevance as usual,




Figure 1: Validity gating (Eq. (3)). An overruling at  $t$  closes Precedent A, so a query at  $t_q > t$  zeroes its indicator and suppresses it despite lexical similarity, while still-open Precedent B is retrieved—the time-valid distinction the RQ3 ablation isolates.

and PLRE receives only support whose legal status is operative at  $t_q$ . This differs from temporal features in a learned reranker, which can still rank closed authority highly, and from citator-style warnings, which are applied after evidence selection.

**Why not a date filter?** Once intervals are known, Eq. (3) is a metadata filter that any vector store can apply, and we keep it deliberately simple. The difficulty lies upstream: the end point  $e_v$  is not an attribute of document  $v$  but is set by *later* documents that overrule, supersede, or narrow it, so it must be inferred from treatment edges extracted from opinion text and propagated through the graph. A filter on document dates cannot express this. Our contribution is therefore the extracted and validated treatment layer—evaluated against gold spans (Table A.5) and a commercial citator (P2), with its error propagation audited (Table 5)—together with evidence that applying it at retrieval time changes PLRE outcomes (Table 4).

### Graph construction and contamination controls.

The graph is built in three stages. *Node induction*: each case, order, or rule from CourtListener/RECAP becomes a node with filing/decision dates from docket metadata and validity interval initialized to  $[s_v, \infty)$ . *Edge extraction*: treatment relations are extracted from opinion text by a two-pass procedure—a citation-sentence detector locates candidate treatment spans, and an LLM classifier (GPT-4o,  $T=0$ , few-shot) labels the relation and polarity. To avoid self-contamination from generated legal conclusions, LLM outputs are treated only as candidate relation labels: interval-closing edges must be anchored to a detected citation span, docket/decision date, and source document identifier, and gold-set evaluation is performed on ci-

tation spans that were not used in prompt examples. *Interval closure*: a confirmed incoming OVERRULES/SUPERSEDES/material-NARROWS edge at date  $t$  closes the target interval ( $e_v \leftarrow t$ ), so retrieval at  $t_q \geq t$  no longer treats it as fully operative. The current graph spans 1,812 nodes and 4,760 treatment edges over the dataset window.

### Edge-extraction accuracy and propagation.

Low-confidence or conflicting negative-treatment edges do not close intervals automatically; they are routed to a human review queue and reports mark the authority as unresolved until review. Because the central claim depends on treatment edges being *correct*, we validate the extractor against a gold set of 600 citation spans hand-labeled by the annotation team (relation type and polarity; adjudicated on disagreement). The negative-treatment relations that drive interval closure (SUPERSEDES, OVERRULES, material NARROWS) are operationally critical—a missed closure is what allows a stale precedent to leak into a risk decision—and attain micro-averaged precision/recall of 0.84/0.77, versus 0.94/0.96 for the high-frequency CITES relation; the full per-relation breakdown is in Table A.5 (appendix). We quantify the downstream effect of missed closures in Table 5 (§6).

### 4.5 Product–Litigation Risk Entailment (PLRE)

We cast product–litigation alignment as a risk-entailment task. Given a product spec  $P$  with attributes  $\{p_1, \dots, p_m\}$  (training-data sources, collection method, output modality, architecture, deployment context, and capability claims), a legal-risk theory  $T_j$ , and a query date  $t_q$ , PLRE predicts:

$$y_{ij} = \mathbf{1}[P, p_i, \mathcal{P}_{t_q} \models T_j]. \quad (4)$$

Table 1 gives schematic instances. Unlike symmetric similarity, PLRE is directional: an attribute such as “web-scale image scraping” can support exposure to a training-data copyright theory, while the inverse relation is not meaningful. Unlike standard legal textual entailment, the premise includes product attributes and the entailment relation is conditioned on time-valid precedent. A pair  $(p_i, T_j)$  is promoted to the risk register only when (i) the entailment score exceeds the validation-set threshold, (ii) the top supporting precedents are valid at  $t_q$ , and (iii) the report includes mitigation text

Table 1: Schematic PLRE instances (illustrative; not drawn from ALERT-Dataset). The same attribute–theory pair can be positive or negative depending on whether its supporting authority is still operative at  $t_q$ .

| Attribute $p_i$                                         | Theory $T_j$                         | Support at $t_q$                                           | $y_{ij}$ |
|---------------------------------------------------------|--------------------------------------|------------------------------------------------------------|----------|
| web-scale image scraping for training                   | training-data copyright infringement | operative authority on unlicensed training copies          | 1        |
| same                                                    | same                                 | only whose interval closed before $t_q$ (stale distractor) | 0        |
| output filter re-moves copyright-management information | DMCA §1202 (output CMI)              | no authority in the query’s jurisdiction                   | 0        |

grounded in at least three cited authorities when available.

PLRE gold labels are built from anonymized product-attribute cards and docket-grounded theory cards. Annotators judge a directional pair as positive only when the product attribute, legal theory, query date, and at least one valid supporting authority jointly support exposure; negative pairs include ordinary mismatches and stale-authority distractors. Crucially, the validity judgment in the labels is made by annotators reading docket and opinion text, *independently of ALERT’s automatically extracted treatment graph*, so the gate must recover label validity rather than presuppose it; an extractor-independent re-check of these validity judgments against an authoritative citator is the downstream link of Protocol P2. Product archetypes are tracked in the split manifest, and year/litigation type are stratified to reduce temporal and topical leakage.

The formulation is system-agnostic: PLRE is directional, context-aware entailment with an exogenous, time-indexed support constraint, and any system—product–theory similarity, legal NLI, temporal reranking, or generic RAG—can be evaluated on it. It is harder than ordinary legal NLI because a model must align engineering terms to legal predicates *and* reject stale but lexically attractive precedent.

## 5 ALERT-Dataset

ALERT-Dataset is a curated benchmark of AI-related copyright and intellectual-property litigation for structured risk analysis.

Table 2: ALERT-Dataset summary statistics. The 3-class risk label merges *critical+high* into *High*. Status counts (initiated/active/settled/dismissed) sum to all 1,247 cases.

| Attribute                                    | Value            |
|----------------------------------------------|------------------|
| Total cases                                  | 1,247            |
| Total documents                              | 8,934            |
| Coverage period                              | 2020–2025        |
| Unique defendants                            | 43               |
| Risk labels                                  | 3 (High/Med/Low) |
| Status dist. (Init/Active/Settled/Dismissed) | 298/695/148/106  |
| Product-mapping GT cases                     | 312              |

### 5.1 Construction and Statistics

ALERT-Dataset compiles CourtListener/RECAP filings and supplementary legal-news signals from January 2020 to December 2025, with filings concentrated in 2023–2025. Cases come from a retrospective keyword search of RECAP complaints, manually screened for AI development, training, or deployment at issue; one case is one docket, so consolidated actions count per docket (App. A.1). Each record follows the 22-field `alert-sensing.csv` schema (Table A.1, appendix). Table 2 summarizes the final annotated subset; raw annual ingestion counts by litigation type are in Table A.2 (appendix).

Each record contains docket metadata, parties, document URLs, and a human-verified litigation summary. Status labels encode workflow state rather than adjudicated outcome; severity is reported separately and is not interpreted as liability.

### 5.2 Annotation and Availability

An automated pre-labeler first assigns a triage score using Table A.3; three legal research specialists then independently review a stratified sample of 400 cases under a fixed guideline. During adjudication, annotators see docket text, product attributes, candidate theory cards, and query dates, but not system identity or the pre-labeler’s score. The remaining 847 cases retain automated pre-labels, with a 5% spot-check showing 91% agreement with at least one specialist (a lenient criterion). Inter-annotator agreement on the 400-case sample is Cohen’s  $\kappa = 0.74$  (Landis and Koch, 1977). Because retained labels can inflate metrics, we release per-record provenance, keep all primary claims on the human-only evaluation, and use retained-label results only as scale-oriented comparisons. The top three defendants account for 41.3%, limiting rep-

representativeness for smaller vendors and non-U.S. jurisdictions. The dataset, guidelines, split files, seeds, evaluation scripts, and logs are provided as anonymized supplementary material for review (App. G).

## 6 Experiments

### 6.1 Experimental Setup

We address six research questions: triage-label agreement (RQ1), report quality (RQ2), loop and temporal-graph ablations (RQ3), schema completeness (RQ4; an operational check reported in Table A.7, appendix), PLRE (RQ5), and evidence-coverage risk control (RQ6).

**Dataset split and label provenance.** We partition ALERT-Dataset into stratified 70/10/20 train/validation/test splits by year and litigation type. Because 847 of 1,247 labels originate from an automated pre-labeler and 400 are human-adjudicated, primary significance claims are computed only on the human-adjudicated subset; retained-label results are reported only for scale. This keeps the rule-based pre-labeler out of the evidence for our primary claims. For ALERT and B2–B5, the training split is used only to select few-shot examples and tune thresholds; no LLM is fine-tuned.

**Baselines.** We compare ALERT with five risk-classification baselines: **B1 Rule-Based** (the ingestion-layer scorer); **B2 Single-Pass LLM**; **B3 RAG-only**; **B4 Single-Agent** (ALERT without the goal-conditioned loop); and **B5 Iteration-Matched Self-Refinement** (a Reflexion-style baseline (Shinn et al., 2023) given ALERT’s mean budget of 2.4 iterations, matched on iterations rather than tokens). B2–B5 use the same GPT-4o backbone and decoding settings as ALERT. For PLRE we add non-ALERT baselines that isolate the suspected confounds: product–theory cosine, LegalBERT product–theory NLI, COLIEE-style (Rabelo et al., 2022) precedent retrieval with no interval closure, and ALERT without temporal gating. These baselines separate product–theory matching, legal entailment, retrieval quality, test-time compute, and the hard validity gate.

**Evaluation metrics.** We report macro F1/precision/recall over three severity labels, report-quality ratings by two blinded specialists, loop efficiency, schema completeness, PLRE P/R/F1, edge-error propagation, and evidence-coverage omission risk under calibrated thresholds. Severity-label F1 measures triage-label agreement,

Table 3: Macro severity-label agreement. F1/P/R are on the human-adjudicated test subset (primary result;  $\pm$  are half-widths of paired-bootstrap confidence intervals, 10,000 resamples). The last column is F1 on the full held-out test set, which includes retained pre-labels (details in Table A.6).

| System             | Human-adjudicated                 |                                   |                                   | Full F1     |
|--------------------|-----------------------------------|-----------------------------------|-----------------------------------|-------------|
|                    | F1                                | P                                 | R                                 |             |
| B1 Rule-Based      | 0.55 $\pm$ 0.03                   | 0.52 $\pm$ 0.04                   | 0.59 $\pm$ 0.04                   | 0.62        |
| B2 Single-Pass LLM | 0.70 $\pm$ 0.03                   | 0.72 $\pm$ 0.03                   | 0.68 $\pm$ 0.03                   | 0.72        |
| B3 RAG-only        | 0.72 $\pm$ 0.03                   | 0.73 $\pm$ 0.03                   | 0.71 $\pm$ 0.03                   | 0.74        |
| B4 Single-Agent    | 0.74 $\pm$ 0.02                   | 0.75 $\pm$ 0.03                   | 0.73 $\pm$ 0.03                   | 0.76        |
| B5 Self-Refinement | 0.75 $\pm$ 0.02                   | 0.76 $\pm$ 0.03                   | 0.74 $\pm$ 0.03                   | 0.77        |
| <b>ALERT Full</b>  | <b>0.79 <math>\pm</math> 0.02</b> | <b>0.80 <math>\pm</math> 0.02</b> | <b>0.78 <math>\pm</math> 0.03</b> | <b>0.81</b> |

not legal liability. Significance uses a paired bootstrap over test predictions (10,000 resamples); calibrated omission risk and backbone robustness are reported in the appendix, with hyperparameters in Table A.10.

### 6.2 RQ1: Triage-Label Agreement

Table 3 reports severity-label agreement on the human-adjudicated subset, our primary circularity check. ALERT outperforms the iteration-matched self-refinement baseline (B5) by 4 F1 points (paired bootstrap,  $p < 0.05$ ). The marginal intervals overlap, but the paired bootstrap scores per-item differences on shared test cases and thus removes the between-case variance that dominates them; we nonetheless read the gap over B5 as modest, especially since the number of human-adjudicated test cases is not reported (about 80 if the 400 are spread evenly across splits). On the full held-out test set (including retained pre-labels; last column and Table A.6) the same ordering holds.

### 6.3 RQ2: Report Quality

Two blinded legal specialists rated the reports of ALERT and B2–B5 for the same 15 cases (75 reports) on 5-point Accuracy, Completeness, Actionability, Evidence Sufficiency, and Clarity scales. ALERT scores 4.3/4.1/4.2/4.4/3.9; case-paired Wilcoxon signed-rank tests give  $p < 0.05$  against every baseline, with the largest gain in Evidence Sufficiency. With 15 pairs per test, no correction for the 20 comparisons, and no inter-rater agreement, RQ2 is supporting evidence only.

### 6.4 RQ3: Agentic Loop and Temporal Graph Ablations

Because temporal precedent reasoning is the central claim, Table 4 decomposes validity-interval gating and treatment-edge weighting. Removing

Table 4: Temporal precedent-graph ablation (main claim). Both mechanisms of Eq. (3) are removed in sequence; the bottom row recovers flat semantic similarity. Stale-cite measures citations to already-closed precedent ( $t_q \geq e_v$ ). Only the full-vs-flat difference was significance-tested ( $p < 0.01$ ).

| Variant                                 | PLRE F1 | $\Delta$ F1 | Stale-cite % |
|-----------------------------------------|---------|-------------|--------------|
| <b>ALERT (Full: gating + weighting)</b> | 0.78    | —           | 4.1          |
| – treatment weighting $w(\cdot)$        | 0.75    | –3          | 7.8          |
| – interval gating (flat similarity)     | 0.72    | –6          | 12.7         |

Table 5: Edge-error propagation audit. “Mask  $x\%$ ” rows randomly remove  $x\%$  of the negative-treatment closures produced by the extractor, on top of its own misses. Effective miss rate relative to gold closures is approximated as  $1 - 0.77(1 - x)$  from the extractor’s 0.77 recall.  $\Delta$ F1 is relative to the extractor row.

| Closure setting         | Eff. miss      | PLRE F1 | $\Delta$ F1 | Stale-cite % |
|-------------------------|----------------|---------|-------------|--------------|
| Gold negative edges     | 0%             | 0.80    | +2          | 3.7          |
| Extractor output        | $\approx 23\%$ | 0.78    | —           | 4.1          |
| + mask 10% of extracted | $\approx 31\%$ | 0.76    | –2          | 6.2          |
| + mask 23% of extracted | $\approx 41\%$ | 0.74    | –4          | 8.9          |
| + mask 30% of extracted | $\approx 46\%$ | 0.72    | –6          | 11.3         |

treatment weighting and then interval gating each degrades PLRE F1. Only the full-vs-flat gap of 6 F1 points was tested for significance ( $p < 0.01$ ); the intermediate 3-point step is reported without an interval and should be read as directional. Stale citations rise from 4.1% to 12.7%. The failure mode is exactly what Eq. (3) targets: flat retrieval can rank a lexically close but already-closed authority above a still-operative one, changing which evidence reaches the generator rather than merely adding a warning after generation. This is the paper’s main methodological result.

Table 5 measures how closure errors propagate. The extractor row already operates at the observed miss rate (about 23% of gold closures, implied by 0.77 recall) and costs 2 PLRE F1 points relative to gold edges. Rows below additionally mask a random fraction of the *extracted* closures before retrieval, simulating a weaker extractor: removing a further 23% of extracted closures ( $\approx 41\%$  effective miss rate) roughly doubles the stale-cite rate (4.1% $\rightarrow$ 8.9%) and lowers PLRE F1 by 4 points. This is why low-confidence negative edges do not close intervals automatically but become unresolved goal atoms (§4.4).

Component ablations for the agentic loop confirm each part contributes: removing the Goal Checker, Planner, and Retriever query-expansion lowers risk-classification F1 from 0.81

Table 6: PLRE comparison including external legal-NLI/precedent-retrieval baselines. External baselines do not use ALERT’s goal predicate or validity-interval closure. Point estimates without confidence intervals.

| System                          | Precision   | Recall      | F1          |
|---------------------------------|-------------|-------------|-------------|
| Product–theory cosine           | 0.64        | 0.66        | 0.65        |
| LegalBERT product–theory NLI    | 0.69        | 0.68        | 0.68        |
| COLIEE-style retrieval, no gate | 0.71        | 0.70        | 0.70        |
| RAG-only generation             | 0.73        | 0.74        | 0.73        |
| ALERT without temporal gating   | 0.72        | 0.73        | 0.72        |
| <b>ALERT Full</b>               | <b>0.77</b> | <b>0.80</b> | <b>0.78</b> |

to 0.76/0.78/0.79 respectively, and only query-expansion removal degrades retrieval (Recall@10 0.84 $\rightarrow$ 0.76); details are in Table A.8.

## 6.5 RQ5: Product–Litigation Risk Entailment

PLRE is evaluated on anonymized AI product specs spanning five archetypes and 312 product-mapping cases, so it is an initial benchmark. We report its scale at the case level; pair-level statistics (number of attribute–theory–date pairs, positive rate, share of stale-authority distractors) and a PLRE-specific inter-annotator agreement are not reported (the  $\kappa$  in §5 is for severity labels). Positive labels require a product attribute, a legal-risk theory, a query date, and time-valid supporting authority; negative labels include ordinary mismatches and stale-authority distractors. In Table 6, ALERT reaches 0.78 F1, 5–13 points above the external baselines; removing temporal gating costs 6 points (decomposed in Table 4). Because positive labels require time-valid support and negatives include stale-authority distractors, the benchmark partly rewards what the gate enforces, and we report no distractor-free comparison; with no confidence intervals and a small benchmark, we read them as directional evidence for the validity-gated formulation, not as coverage of all AI-product litigation.

## 6.6 RQ6: Evidence-Coverage Risk Control and Temporal Shift

Calibrating the evidence threshold to  $\alpha \in \{0.05, 0.10, 0.20\}$  yields omission risk 0.043/0.082/0.171 in random splits. Under a temporal split (2023–24  $\rightarrow$  2025) the achieved risk rises to 0.061/0.113/0.196, exceeding the target at  $\alpha=0.05$  and 0.10 (Table A.11), which motivates rolling re-calibration. To probe dependence on GPT-4o, we re-run the full system with an open-weights Llama-3.1-70B+bge stack: it runs unchanged, with macro F1 dropping from

0.81 to 0.76 and similar loop length (2.4 vs. 2.6 iterations; Table A.12). We did not re-run the baselines on this stack, so this shows only that the pipeline transfers to an open-weights backbone, not that its advantage over baselines is backbone-independent.

## 7 Related Work

**Legal NLP and entailment.** Legal NLP has studied judgment prediction (Chalkidis et al., 2019), statutory reasoning (Louis and Spanakis, 2023), legal understanding benchmarks (Chalkidis et al., 2022), case retrieval and entailment (Rabelo et al., 2022), and summarization (Huang et al., 2021; Manor and Li, 2019), often with domain-adapted encoders (Chalkidis et al., 2020; Zheng et al., 2021; Xiao et al., 2021) built on pretrained transformers (Vaswani et al., 2017; Devlin et al., 2019). LegalBench evaluates legal reasoning (Guha et al., 2023), and audits find that LLMs and commercial legal RAG tools still hallucinate (Dahl et al., 2024; Magesh et al., 2024). These works take legal text as the premise, treat the corpus as static, and do not test whether retrieved authority was still operative at a given date. PLRE changes the inference unit: it maps product attributes to legal-risk theories, and the relation is directional, product-conditioned, and temporal.

**Time-sensitive QA and temporal retrieval.** A line of work shows that answers depend on when a question is asked: SituatedQA conditions QA on temporal and geographic context (Zhang and Choi, 2021), TimeQA and StreamingQA evaluate time-sensitive and evolving knowledge (Chen et al., 2021; Liška et al., 2022), RealTime QA targets answers that change over time (Kasai et al., 2023), and time-aware LMs condition parametric knowledge on time (Dhingra et al., 2022); temporal KG QA reasons over time-stamped facts (Saxena et al., 2021), and temporal IR uses document time as a ranking signal (Campos et al., 2014). These settings treat time as a property of *facts* or *queries*. In PLRE, time enters through *supersession*: a later document changes the evidential status of an earlier one, so validity must be propagated along treatment edges rather than read from a timestamp.

**Citators and legal citation graphs.** Commercial citators such as Shepard’s and KeyCite track whether authority has been overruled, questioned, or negatively treated (LexisNexis, 2024; Thomson Reuters, 2024); we do not claim time-varying

authority itself as new. Citation-network analysis measures the importance of precedent (Fowler et al., 2007), and prior work labels the semantics of citation edges automatically (Sadeghian et al., 2018). These efforts build or characterize the graph; ALERT instead attaches validity intervals to extracted treatment edges and applies them as a *retrieval-time gate* (Eq. (3), Fig. 1) that feeds a downstream product-risk task.

**RAG, agents, and risk control.** RAG (Lewis et al., 2020; Karpukhin et al., 2020; Gao et al., 2023) usually runs a fixed retrieval plan before generation, while agent frameworks (Yao et al., 2023; Shinn et al., 2023; Wei et al., 2022; LangChain, Inc., 2024; Gravitas, 2023; Park et al., 2023; Wang et al., 2024a,b) show gains from iterative tool use; self-refinement, however, stops when the model judges its own answer satisfactory. ALERT ties stopping to an externally inspectable predicate over schema completeness, evidence support, and mitigation coverage, and uses conformal risk control (Angelopoulos et al., 2023, 2021; Geifman and El-Yaniv, 2017) narrowly for evidence omission, not for legal correctness.

## 8 Discussion

**What drives the gains.** The gains come from knowing *which* authority is still operative—the extracted, validated treatment layer—and from applying that knowledge before generation; the gate itself (Eq. (3)) is intentionally simple, and the agents and ingestion schema are supporting infrastructure. Table 4 isolates the effect of applying the layer, and B5 rules out extra refinement passes, though not extra LLM calls, as the source of the gain.

**Construct validity.** Agreement of the treatment graph with a commercial citator is measured (P2, Appendix I.2): negative-treatment micro-F1 0.81, 20.4% missed closures, and a 3.8% PLRE decision-flip rate after citator correction. Transfer of the gate to external retrieval stacks (P1, Appendix I.1) is specified in advance but not yet executed.

## 9 Conclusion

We introduced PLRE, a directional, time-conditioned entailment task linking AI-product attributes to legal-risk theories, and ALERT, whose extracted, citator-validated treatment layer tells retrieval which authority was operative at the query date. Validity gating improves PLRE F1 and cuts stale citations relative to flat similarity.

## Limitations

**Guarantees.** Proposition 1 is a bounded termination-and-escalation property, not a convergence theorem. Its termination clause is close to definitional, and its progress assumption is not verified directly. Calibration controls evidence omission under its calibration distribution (whose size we do not report), not correct legal conclusions, temporal-shift robustness, or freedom from LLM hallucinations (Dahl et al., 2024). Because litigation streams are non-stationary, the omission guarantee is weakest under the temporal split (Table A.11)—precisely where deployment needs it most—so we treat rolling or weighted recalibration as a deployment requirement rather than an optional refinement.

**Generality is untested.** The validity-gating operator is defined generally, but all evidence comes from ALERT’s own pipeline on U.S. AI-litigation data. Its transfer to *external* retrieval stacks and corpora is specified only as a protocol (Protocol P1, Appendix I.1) and has not been executed, so the gains reported here should not be read as evidence that the gate improves other retrievers or other domains. Likewise, backbone robustness is limited: the open-weights re-run (Table A.12) covers only the full ALERT system, not the baselines, so we cannot yet claim that ALERT’s margin over baselines holds independently of the LLM backbone. Cross-stack transfer and backbone-matched baseline comparisons therefore remain the main open generality tests.

**Benchmark scale and evaluation.** At the current scale PLRE remains an initial benchmark over five product archetypes and 312 product-mapping cases. We do not report its pair-level composition (positive rate, share of stale-authority distractors), a PLRE-specific inter-annotator agreement, confidence intervals for Tables 4 and 6 beyond the full-vs-flat test, or results on a subset without stale-authority distractors. Since PLRE labels encode temporal validity, part of the gating gain may come from rejecting such distractors rather than from better entailment. PLRE results should therefore be read as directional evidence for the validity-gated formulation rather than as precise effect sizes or coverage of all AI-product litigation. The report-quality study (RQ2) uses two blinded raters on 15 cases, without inter-rater agreement or multiple-comparison correction, and is treated as supporting, not primary, evidence. Because 847 of 719,247 la-

bels are retained pre-labels, all primary significance claims are restricted to the human-adjudicated subset, whose test-split size we do not report, and effect sizes over the strongest baseline are modest (Table 3).

**Coverage and jurisdiction.** Cases are counted per docket, so disputes litigated across several related or consolidated dockets are over-represented relative to distinct disputes (App. A.1). The dataset is also skewed toward the largest defendants—the top three account for 41.3% of cases—limiting representativeness for smaller vendors. ALERT targets U.S. federal CourtListener/RECAP coverage and English-language documents; non-U.S. jurisdictions require different docket schemas and citation practices, and extending the treatment graph to them is future work.

**Extractor recall.** Negative-treatment edge recall (0.77 internal gold; 20.4% missed-closure rate against an authoritative citator) is useful but not lossless, and a missed closure is exactly what allows a stale precedent to leak into a risk decision; ALERT mitigates this with human-review routing for low-confidence closures, but improving extraction recall remains an open problem.

## Ethics Statement

**Decision-support boundary.** ALERT emits risk intelligence, evidence summaries, and mitigation candidates with provenance and residual uncertainty for qualified human reviewers. It is not legal advice: it does not replace counsel, predict litigation outcomes, or recommend business decisions, and severity labels are triage proxies that are not interpreted as liability. Every High- or Medium-risk item is presented for review by a qualified person, and unresolved legal-status atoms are surfaced as such rather than resolved automatically.

**Human oversight of failure modes.** Low-confidence negative-treatment edges enter a human review queue, since missed closures are the failure mode most likely to admit stale authority into a report. Deployment logs store retrieved authorities, validity intervals, and unmet goal atoms so reviewers can distinguish a genuine product-theory match from a retrieval gap or an unresolved legal-status ambiguity.

**Data governance.** Released data are public docket metadata and machine-generated summaries; we do not redistribute attorney-authored complaint text beyond what is already public,

and the P2 citator study releases only the cross-walk, sampling seeds, and agreement statistics, with no redistributed proprietary citator content. Source records are public U.S. court filings obtained through CourtListener/RECAP and are used for research consistent with that service’s terms; news items are represented only by links and derived metadata. Released records contain the names of parties and counsel as they appear in these public filings; we collect no other personal information. ALERT-Dataset is released under CC BY 4.0 and the code under the MIT license, for research on legal-risk analysis; neither is intended for automated legal decisions.

**Annotators and raters.** The three legal research specialists who adjudicated labels and the two specialists who rated reports (§5, §6) are staff of the authors’ institution and performed this work within their regular paid duties, without additional compensation. They were told how their annotations would be used and agreed verbally. Because the work consisted of labeling public court documents and collected no personal information about the participants, we did not seek formal ethics-board review. Annotation guidelines are released with the dataset.

**Use of AI assistants.** AI writing assistants were used only to edit and polish the text of this paper. Research ideas, system design, experiments, and analyses are the authors’ own.

## References

- Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. 2021. [Learn then test: Calibrating predictive algorithms to achieve risk control](#). arXiv preprint arXiv:2110.01052.
- Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. 2023. [Conformal risk control](#). arXiv preprint arXiv:2208.02814.
- Ricardo Campos, Gaël Dias, Alípio M. Jorge, and Adam Jatowt. 2014. Survey of temporal information retrieval and related applications. *ACM Computing Surveys*, 47(2):15:1–15:41.
- Ilias Chalkidis, Ion Androutsopoulos, and Nikolaos Aletras. 2019. [Neural legal judgment prediction in English](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4317–4323.
- Ilias Chalkidis, Manos Fergadiotis, Prodromos Malakasiotis, Nikolaos Aletras, and Ion Androutsopoulos. 2020. [LEGAL-BERT: The muppets straight out of law school](#). In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2898–2904. Association for Computational Linguistics.
- Ilias Chalkidis, Abhik Jana, Dirk Hartung, Michael Bommarito, Ion Androutsopoulos, Daniel Katz, and Nikolaos Aletras. 2022. [LexGLUE: A benchmark dataset for legal language understanding in English](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4310–4330.
- Wenhu Chen, Xinyi Wang, and William Yang Wang. 2021. A dataset for answering time-sensitive questions. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*.
- Matthew Dahl, Varun Magesh, Mirac Suzgun, and Daniel E. Ho. 2024. [Large legal fictions: Profiling legal hallucinations in large language models](#). *Journal of Legal Analysis*, 16(1):64–93.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*, pages 4171–4186. Association for Computational Linguistics.
- Bhuwan Dhingra, Jeremy R. Cole, Julian Martin Eisenschlos, Daniel Gillick, Jacob Eisenstein, and William W. Cohen. 2022. [Time-aware language models as temporal knowledge bases](#). *Transactions of the Association for Computational Linguistics*, 10:257–273.
- James H. Fowler, Timothy R. Johnson, James F. Spriggs, Sangick Jeon, and Paul J. Wahlbeck. 2007. Network analysis and the law: Measuring the legal importance of precedents at the U.S. Supreme Court. *Political Analysis*, 15(3):324–346.
- Free Law Project. 2023. [CourtListener: A free legal research tool](#). Online platform and REST API v4.
- Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Qianyu Guo, Meng Wang, and Haofen Wang. 2023. [Retrieval-augmented generation for large language models: A survey](#). Preprint, arXiv:2312.10997.
- Yonatan Geifman and Ran El-Yaniv. 2017. [Selective classification for deep neural networks](#). In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30.
- Getty Images, Inc. 2023. [Getty images, inc. v. stability AI, inc., case no. 1:23-cv-00135](#). U.S. District Court, D. Delaware.
- Significant Gravitas. 2023. [AutoGPT: An autonomous GPT-4 experiment](#). Open-source software.
- James Grimmelman. 2016. [Copyright for literate robots](#). *Iowa Law Review*, 101:657–681.

|     |                                                                                |                                                                               |     |
|-----|--------------------------------------------------------------------------------|-------------------------------------------------------------------------------|-----|
| 886 | Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher                            | knowledge over time in question answering models.                             | 941 |
| 887 | Ré, et al. 2023. LegalBench: A collaboratively built                           | In <i>Proceedings of the 39th International Conference</i>                    | 942 |
| 888 | benchmark for measuring legal reasoning in large                               | on <i>Machine Learning</i> , volume 162 of <i>Proceedings of</i>              | 943 |
| 889 | language models. In <i>Advances in Neural Information</i>                      | <i>Machine Learning Research</i> .                                            | 944 |
| 890 | <i>Processing Systems (Datasets and Benchmarks</i>                             |                                                                               |     |
| 891 | <i>Track</i> ).                                                                |                                                                               |     |
| 892 | Peter Henderson, Xuechen Li, Mark Lemley, Dan Ju-                              | Anne Louis and Gerasimos Spanakis. 2023. <a href="#">Statutory</a>            | 945 |
| 893 | rafsky, and Daniel E. Ho. 2023. <a href="#">Foundation models</a>              | <a href="#">reasoning assessment with a faithful explanations</a>             | 946 |
| 894 | <a href="#">and fair use</a> . <i>Journal of Machine Learning Research</i> ,   | <a href="#">dataset</a> . In <i>Proceedings of the 17th Conference of the</i> | 947 |
| 895 | 24(400):1–79.                                                                  | <i>European Chapter of the Association for Computa-</i>                       | 948 |
|     |                                                                                | <i>tional Linguistics (EACL)</i> , pages 979–994. Associa-                    | 949 |
|     |                                                                                | tion for Computational Linguistics.                                           | 950 |
| 896 | Luyang Huang, Shuyang Cao, Nikolaus Parulian, Heng                             | Varun Magesh, Faiz Surani, Matthew Dahl, Mirac Suz-                           | 951 |
| 897 | Ji, and Lu Wang. 2021. <a href="#">Efficient attentions for long</a>           | gun, Christopher D. Manning, and Daniel E. Ho.                                | 952 |
| 898 | <a href="#">document summarization</a> . In <i>Proceedings of the 2021</i>     | 2024. Hallucination-free? assessing the reliability                           | 953 |
| 899 | <i>Conference of the North American Chapter of the</i>                         | of leading AI legal research tools. arXiv preprint                            | 954 |
| 900 | <i>Association for Computational Linguistics (NAACL-</i>                       | arXiv:2405.20362.                                                             | 955 |
| 901 | <i>HLT</i> ), pages 1419–1428. Association for Computa-                        |                                                                               |     |
| 902 | tional Linguistics.                                                            | Laura Manor and Junyi Jessy Li. 2019. <a href="#">Plain English</a>           | 956 |
| 903 | Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick                            | <a href="#">summarization of contracts</a> . In <i>Proceedings of the</i>     | 957 |
| 904 | Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and                               | <i>Natural Legal Language Processing Workshop 2019</i> ,                      | 958 |
| 905 | Wen-tau Yih. 2020. <a href="#">Dense passage retrieval for open-</a>           | pages 1–11. Association for Computational Linguis-                            | 959 |
| 906 | <a href="#">domain question answering</a> . In <i>Proceedings of the</i>       | tics.                                                                         | 960 |
| 907 | <i>2020 Conference on Empirical Methods in Natural</i>                         | Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai,                             | 961 |
| 908 | <i>Language Processing (EMNLP)</i> , pages 6769–6781.                          | Meredith Ringel Morris, Percy Liang, and Michael S.                           | 962 |
| 909 | Association for Computational Linguistics.                                     | Bernstein. 2023. <a href="#">Generative agents: Interactive simu-</a>         | 963 |
| 910 | Jungo Kasai, Keisuke Sakaguchi, Yoichi Takahashi,                              | <a href="#">lacra of human behavior</a> . In <i>Proceedings of the 36th</i>   | 964 |
| 911 | Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir                                 | <i>Annual ACM Symposium on User Interface Software</i>                        | 965 |
| 912 | Radev, Noah A. Smith, Yejin Choi, and Kentaro Inui.                            | <i>and Technology (UIST)</i> , pages 1–22. ACM.                               | 966 |
| 913 | 2023. RealTime QA: What’s the answer right now?                                | Juliano Rabelo, Randy Goebel, Mi-Young Kim, Yoshi-                            | 967 |
| 914 | In <i>Advances in Neural Information Processing Sys-</i>                       | nobu Kano, Masaharu Yoshioka, and Ken Satoh.                                  | 968 |
| 915 | <i>tems (Datasets and Benchmarks Track)</i> .                                  | 2022. Overview and discussion of the competi-                                 | 969 |
| 916 | Sarah Kosari and Meta Platforms, Inc. 2024. <a href="#">Kosari et</a>          | tion on legal information extraction/entailment (COL-                         | 970 |
| 917 | <a href="#">al. v. meta platforms, inc., case no. 3:24-cv-04722</a> .          | IEE) 2021. <i>The Review of Socionetwork Strategies</i> ,                     | 971 |
| 918 | U.S. District Court, N.D. California.                                          | 16(1):111–133.                                                                | 972 |
| 919 | J. Richard Landis and Gary G. Koch. 1977. <a href="#">The mea-</a>             | Ali Sadeghian, Lakshman Sundaram, Daisy Zhe Wang,                             | 973 |
| 920 | <a href="#">surement of observer agreement for categorical data</a> .          | William F. Hamilton, Karl Branting, and Craig                                 | 974 |
| 921 | <i>Biometrics</i> , 33(1):159–174.                                             | Pfeifer. 2018. Automatic semantic edge labeling                               | 975 |
| 922 | LangChain, Inc. 2024. <a href="#">LangGraph: Building stateful,</a>            | over legal citation graphs. <i>Artificial Intelligence and</i>                | 976 |
| 923 | <a href="#">multi-actor applications with LLMs</a> . Open-source               | <i>Law</i> , 26(2):127–144.                                                   | 977 |
| 924 | software.                                                                      | Apoorv Saxena, Soumen Chakrabarti, and Partha Taluk-                          | 978 |
| 925 | Amanda Levendowski. 2018. <a href="#">How copyright law can fix</a>            | dar. 2021. <a href="#">Question answering over temporal knowl-</a>            | 979 |
| 926 | <a href="#">artificial intelligence’s implicit bias problem</a> . <i>Wash-</i> | <a href="#">edge graphs</a> . In <i>Proceedings of the 59th Annual</i>        | 980 |
| 927 | <i>ington Law Review</i> , 93(2):579–630.                                      | <i>Meeting of the Association for Computational Lin-</i>                      | 981 |
| 928 | Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio                           | <i>guistics and the 11th International Joint Conference</i>                   | 982 |
| 929 | Petroni, Vladimir Karpukhin, Naman Goyal, Hein-                                | <i>on Natural Language Processing (Volume 1: Long</i>                         | 983 |
| 930 | rich Küttler, Mike Lewis, Wen-tau Yih, Tim Rock-                               | <i>Papers)</i> , pages 6663–6676.                                             | 984 |
| 931 | täschel, Sebastian Riedel, and Douwe Kiela. 2020.                              | Noah Shinn, Federico Cassano, Ashwin Gopinath,                                | 985 |
| 932 | <a href="#">Retrieval-augmented generation for knowledge-</a>                  | Karthik Narasimhan, and Shunyu Yao. 2023. <a href="#">Re-</a>                 | 986 |
| 933 | <a href="#">intensive NLP tasks</a> . In <i>Advances in Neural Informa-</i>    | <a href="#">flexion: Language agents with verbal reinforcement</a>            | 987 |
| 934 | <i>tion Processing Systems (NeurIPS)</i> , volume 33,                          | <a href="#">learning</a> . In <i>Advances in Neural Information Pro-</i>      | 988 |
| 935 | pages 9459–9474.                                                               | <i>cessing Systems (NeurIPS)</i> , volume 36.                                 | 989 |
| 936 | LexisNexis. 2024. <a href="#">Shepard’s Citation Service</a> . Online          | The New York Times Company. 2023. <a href="#">The new york</a>                | 990 |
| 937 | legal research service documentation.                                          | <a href="#">times company v. microsoft corporation et al., case</a>           | 991 |
| 938 | Adam Liška, Tomáš Kočiský, Elena Gribovskaya, Tay-                             | <a href="#">no. 1:23-cv-11195</a> . U.S. District Court, S.D.N.Y.             | 992 |
| 939 | fun Terzi, Eren Sezener, Devang Agrawal, et al. 2022.                          | Thomson Reuters. 2024. <a href="#">KeyCite on Westlaw Precision</a> .         | 993 |
|     | StreamingQA: A benchmark for adaptation to new                                 | Online legal research service documentation.                                  | 994 |

- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. [Attention is all you need](#). In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 30.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. 2024a. [Voyager: An open-ended embodied agent with large language models](#). *Transactions on Machine Learning Research*.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. 2024b. [A survey on large language model based autonomous agents](#).
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. 2022. [Chain-of-thought prompting elicits reasoning in large language models](#). In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 24824–24837.
- Chaojun Xiao, Xueyu Hu, Zhiyuan Liu, Cunchao Tu, and Maosong Sun. 2021. [Lawformer: A pre-trained language model for Chinese legal long documents](#). *AI Open*, 2:79–84.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. [ReAct: Synergizing reasoning and acting in language models](#). In *International Conference on Learning Representations (ICLR)*.
- Michael Zhang and Eunsol Choi. 2021. [SituatingQA: Incorporating extra-linguistic contexts into QA](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7371–7387.
- Lucia Zheng, Neel Guha, Brandon R. Anderson, Peter Henderson, and Daniel E. Ho. 2021. [When does pre-training help? assessing self-supervised learning for law and the CaseHOLD dataset](#). In *Proceedings of the Eighteenth International Conference on Artificial Intelligence and Law (ICAIL)*, pages 159–168. ACM.

## Technical Appendix

This appendix supplements the main text with material that supports reproducibility but exceeds the page budget of the body. It contains the full 22-field schema, the ingestion-layer scoring matrix, formatted History/Progress Notes examples, open-source implementation details and configuration parameters of alert-tracker, treatment-edge extraction accuracy, full-test-set RQ1 results and the schema-completeness/agent-loop/retriever ablations, goal-predicate hyperparameters with agent prompt templates and evaluation-protocol notes, and the proof of Proposition 1 with the conformal calibration and temporal-shift audit.

### A Full alert-sensing.csv Schema

Table A.1 lists all 22 fields of alert-sensing.csv, with type, requirement level (R = Required, Rec = Recommended, O = Optional), and description.

Table A.1: The full 22-field alert-sensing.csv schema.

| #  | Field           | Type     | Req | Description                 |
|----|-----------------|----------|-----|-----------------------------|
| 1  | Case No.        | Int      | R   | Auto-incrementing index     |
| 2  | Case Title      | String   | R   | Plaintiff v. Defendant      |
| 3  | Filing Date     | Date     | R   | Date filed (YYYY-MM-DD)     |
| 4  | Docket No.      | String   | R   | Court-assigned number       |
| 5  | Plaintiff       | String   | R   | Name(s), comma-separated    |
| 6  | Defendant       | String   | R   | Name(s), comma-separated    |
| 7  | Target Data     | String   | R   | Training data type alleged  |
| 8  | Target Product  | String   | R   | AI product implicated       |
| 9  | Jurisdiction    | String   | R   | Country of filing           |
| 10 | Court           | String   | R   | Official court name         |
| 11 | Status          | Enum     | R   | Lifecycle code              |
| 12 | Last Update     | Date     | R   | Most recent refresh date    |
| 13 | Cause of Action | String   | R   | Core legal claim(s)         |
| 14 | Claim (USD)     | String   | Rec | Damages or statutory range  |
| 15 | Summary         | Long str | Rec | AI-generated case narrative |
| 16 | Tracker URL     | URL      | Rec | CourtListener/PACER link    |
| 17 | Source URL      | String   | O   | Initial discovery source    |
| 18 | History         | Long str | R   | Chronological update log    |
| 19 | Plaintiff Atty. | String   | O   | Lead plaintiff counsel      |
| 20 | Progress Notes  | Long str | Rec | Markdown update log         |
| 21 | Defense Atty.   | String   | O   | Lead defense counsel        |
| 22 | Remarks         | String   | O   | Press links, extra context  |

#### A.1 Inclusion Criteria and Case Unit

**Search.** Candidate dockets were retrieved from CourtListener/RECAP for January 2020–December 2025 by searching complaint, amended-complaint, and petition entries for combinations of three keyword families: (i) AI-training terms (e.g., AI training, model training, training data, dataset, LLM); (ii) legal-theory terms (e.g., copyright, DMCA, unauthorized, without permission, pirated, scraping); and (iii) source-specific terms

for known data sources (e.g., shadow libraries, video datasets, music lyrics). Legal-news signals were used only to locate dockets, not as cases.

**Screening.** A reviewer kept a docket only if the complaint puts the development, training, or deployment of an AI system at issue; keyword matches about unrelated software or data disputes were removed. Litigation type (Table A.2) is assigned afterwards from the recorded cause of action, which is why non-copyright categories (trade secret, other) appear.

**Case unit.** One case corresponds to one docket. Deduplication removes repeated retrievals of the same docket (1,260 raw records to 1,247 cases) but does not merge related or consolidated dockets, so a single dispute litigated in several dockets contributes several cases. Counts should therefore be read as dockets, not as distinct disputes.

Table A.2 gives the raw annual ingestion counts by litigation type, before deduplication and quality screening.

Table A.2: Annual filing breakdown by litigation type. Counts are raw ingested filings (1,260) before deduplication; the annotated subset contains 1,247 cases.

| Year         | Copyright  | DMCA       | Trade Secret | Other     | Total        |
|--------------|------------|------------|--------------|-----------|--------------|
| 2020         | 28         | 9          | 7            | 3         | 47           |
| 2021         | 41         | 14         | 10           | 4         | 69           |
| 2022         | 89         | 31         | 22           | 8         | 150          |
| 2023         | 198        | 72         | 51           | 19        | 340          |
| 2024         | 261        | 96         | 64           | 24        | 445          |
| 2025         | 122        | 47         | 29           | 11        | 209          |
| <b>Total</b> | <b>739</b> | <b>269</b> | <b>183</b>   | <b>69</b> | <b>1,260</b> |

### B Ingestion-Layer Scoring Matrix

Table A.3 lists the additive ingestion-layer scoring matrix (main text) used for real-time triage and as the automated pre-labeler. The critical-composite bonus fires only when copyright and model-training evidence co-occur. This heuristic is a baseline, never ground truth.

### C Formatted Examples: History and Progress Notes

The History field (field 18) records a reverse-chronological timeline using the format YYYY-MM-DD, [event]:

Table A.3: Ingestion-layer scoring matrix. The critical-composite bonus fires only when copyright and model-training evidence co-occur.

| Evidence Category             | Key Indicators                                         | Score      |
|-------------------------------|--------------------------------------------------------|------------|
| <b>Critical composite</b>     | <b>Copyright <math>\wedge</math> Training co-cited</b> | <b>+50</b> |
| Copyright direct mention      | NOS 820, copyright keyword                             | +30        |
| Unauthorized collection       | scrape, crawl, ingest, harvest                         | +25        |
| Direct model training mention | train, LLM, transformer, diffusion                     | +20        |
| Copyright / IP dispute        | infringement, DMCA, fair use                           | +10        |
| Commercial exploitation       | profit, revenue, subscription                          | +10        |
| Class action filing           | class action, putative class                           | +5         |
| Licensed data cooperation     | contract, licensing, partnership                       | -10        |

2025-03-06, Plaintiff requests \$15,000 per infringed work.  
2025-03-01, Complaint filed alleging unauthorized use of  
    copyrighted dataset for AI model training.

The Progress Notes field (field 20) uses Markdown ## headers to distinguish initial from subsequent updates (the entry below is a *synthetic format illustration*, not a real case):

## (Started), 2025-03-17, [Plaintiff] v. [Defendant] --  
    Unauthorized training-data use, commercial use claimed.  
## (Updated), 2025-03-18, Preliminary injunction  
    motion filed by plaintiff.

## D Open-Source Implementation Details

### D.1 System Architecture and Discovery

Figure A.1 shows how the sensing pipeline feeds the agentic layer evaluated in the main text. ALERT’s discovery layer ingests CourtListener/RECAP (Free Law Project, 2023), docket-backed legal-media signals, and keyword-filtered news. A two-stage deduplication module first matches docket IDs and titles, then optionally removes Gemini-embedding neighbors above cosine 0.85 (Algorithm 1, appendix). Detected filings populate the 22-field alert-sensing.csv schema by combining docket metadata, NER over complaint text, regex extraction for damages, and LLM summarization for the narrative fields. All code will be released, as described in Appendix G.




Figure A.1: ALERT architecture. The sensing pipeline (left) detects filings, collects documents, and populates the schema; the agentic layer (right) plans, retrieves time-valid precedent, scores risk, and reports, with the Goal Checker driving a refine loop until evidence coverage is met.

### D.2 Case Study and Live Monitoring

Case studies illustrate mechanism behavior, not outcome prediction. In a held-out PLRE query about output-CMI risk, flat RAG ranked a lexically close authority whose interval had already closed; validity gating removed it from support and promoted a still-operative authority, preventing a stale-citation report. In 30 days, alert-tracker processed 1,159 items, surfaced 94 unique filings after 92% deduplication, and reached discovery-layer Precision 0.90/Recall 0.84.

### D.3 Sensing Tool

The alert-tracker tool runs on an adaptive, configurable GitHub Actions schedule: hourly during a configurable business-hours window and two-hourly off-hours. On-demand execution via workflow\_dispatch is also supported. Its principal modules are:

- **Query engine:** five CourtListener queries (complaint/petition filters with AI keywords, REST API v4) plus four news queries over Google News RSS.
- **Ingestion pipeline:** orchestrates polling, RSS monitoring, docket backtracking, and cross-issue deduplication over a configurable window.
- **Complaint parser:** extracts cause-of-action language, training snippets, and parties from RECAP PDFs.

- **Risk scorer:** applies the seven-category additive matrix *plus* the critical-composite (+50) co-occurrence rule of Table A.3 to produce a 0–100 score with a four-band severity indicator.

- **Hybrid deduplication:** a two-stage pipeline (Algorithm 1) that combines lexical title/docket matching with optional semantic similarity using Gemini text-embedding-004 embeddings (cosine threshold 0.85), enabled by a configuration flag.

- **Multi-modal reporting:** a Gemini Flash trend summary at  $\approx 1\text{--}3$  calls per run, plus an optional Imagen-3 visual brief attached to the day-close summary. A provider-agnostic interface accommodates alternative LLM/diffusion back-ends.

- **Reporting integrations:** posts to GitHub Issues with a three-tier hierarchical structure (issue body / morning trend summary / evening deduplicated table) and a parallel Slack notification.

#### Algorithm 1 Hybrid Lexical–Semantic Deduplication

**Require:** Candidate items  $\mathcal{N}_{\text{new}}$ , baseline items  $\mathcal{N}_{\text{base}}$  (prior issues), threshold  $\sigma$

**Ensure:** Deduplicated set  $\mathcal{N}^*$

- 1:  $T_b \leftarrow \text{EXTRACTTITLES}(\mathcal{N}_{\text{base}})$
- 2:  $\mathcal{R} \leftarrow \{i \in \mathcal{N}_{\text{new}} \mid \text{title}(i) \notin T_b\}$  {Stage 1: exact-title lexical filter}
- 3: **if**  $\text{LLMENABLED} \wedge |\mathcal{R}| > 0$  **then**
- 4:  $E_{\text{curr}} \leftarrow \text{GEMINIEMBED}(\{\text{title}(i) : i \in \mathcal{R}\})$
- 5:  $E_{\text{base}} \leftarrow \text{GEMINIEMBED}(T_b)$
- 6:  $\mathcal{S} \leftarrow \{i \in \mathcal{R} \mid \exists e \in E_{\text{base}} : \cos(E_{\text{curr}}[i], e) \geq \sigma\}$
- 7:  $\mathcal{R} \leftarrow \mathcal{R} \setminus \mathcal{S}$  {Stage 2: semantic filter}
- 8: **end if**
- 9: **return**  $\mathcal{N}^* \leftarrow \mathcal{R}$   $\{\approx 2$  embedding calls/run; free-tier compatible}

Table A.4 lists the key configuration parameters of `alert-tracker`, covering API credentials, scheduling, deduplication, and reporting options. All secrets are injected via GitHub Secrets; environment variables are set via GitHub Variables.

Table A.4: Key configuration parameters of `alert-tracker` (identifiers anonymized for review). Secrets via GitHub Secrets; variables via GitHub Variables.

| Parameter                | Default             | Description                         |
|--------------------------|---------------------|-------------------------------------|
| Lookback window          | 3 days              | Days of history to retrieve         |
| CourtListener token      | —                   | CourtListener API v4 token          |
| Issue title base         | —                   | Base title for GitHub Issues        |
| Issue label              | (project label)     | GitHub Issue label                  |
| LLM API key              | —                   | LLM API key (optional)              |
| LLM model                | gemini-flash-latest | Gemini model id                     |
| Trend window             | (unset)             | Trend-summary window (days)         |
| Semantic dedup           | off                 | Enable semantic dedup (stage 2)     |
| Semantic dedup threshold | 0.85                | Cosine threshold for semantic dedup |
| Daily visual brief       | off                 | Enable Imagen-3 daily visual brief  |
| Cross-issue dedup window | (unset)             | Cross-issue dedup window (days)     |
| Show docket candidates   | off                 | Show uncertain candidates           |
| Collapse long cells      | off                 | Fold long table cells               |
| Collapse article URLs    | off                 | Collapse article URL lists          |
| Debug                    | off                 | Verbose logging                     |

## E Treatment-Edge Extraction Accuracy

Table A.5 reports the per-relation accuracy of the treatment-edge extractor (main text) against the hand-labeled gold span set. Negative-treatment edges (SUPERSEDES/OVERRULES/material NARROWS) drive validity-interval closure and are reported separately from the high-frequency CITES relation, with a micro-average over the three negative-treatment relations.

Table A.5: Treatment-edge extraction accuracy against the hand-labeled gold span set (main text).

| Relation                          | Precision | Recall | F1   |
|-----------------------------------|-----------|--------|------|
| CITES                             | 0.94      | 0.96   | 0.95 |
| DISTINGUISHES                     | 0.81      | 0.76   | 0.78 |
| NARROWS (material)                | 0.79      | 0.71   | 0.75 |
| SUPERSEDES                        | 0.85      | 0.78   | 0.81 |
| OVERRULES                         | 0.88      | 0.82   | 0.85 |
| <b>Negative-treatment (micro)</b> | 0.84      | 0.77   | 0.80 |

## F Additional Results and Ablations

Table A.6: Severity-label agreement on the full held-out test set, including retained pre-labels (RQ1; reported for scale). B5 is matched to ALERT’s mean iteration budget (not to total LLM calls or tokens).

| System             | F1                                | Precision                         | Recall                            |
|--------------------|-----------------------------------|-----------------------------------|-----------------------------------|
| B1 Rule-Based      | 0.62 $\pm$ 0.02                   | 0.58 $\pm$ 0.03                   | 0.67 $\pm$ 0.02                   |
| B2 Single-Pass LLM | 0.72 $\pm$ 0.02                   | 0.74 $\pm$ 0.02                   | 0.70 $\pm$ 0.03                   |
| B3 RAG-only        | 0.74 $\pm$ 0.02                   | 0.75 $\pm$ 0.02                   | 0.72 $\pm$ 0.02                   |
| B4 Single-Agent    | 0.76 $\pm$ 0.02                   | 0.77 $\pm$ 0.02                   | 0.75 $\pm$ 0.02                   |
| B5 Self-Refinement | 0.77 $\pm$ 0.02                   | 0.78 $\pm$ 0.02                   | 0.76 $\pm$ 0.02                   |
| <b>ALERT Full</b>  | <b>0.81 <math>\pm</math> 0.01</b> | <b>0.82 <math>\pm</math> 0.02</b> | <b>0.80 <math>\pm</math> 0.02</b> |

Table A.6 reports severity-label agreement on the full held-out test set (RQ1); primary claims use the human-adjudicated subset (Table 3).

RQ4 is an operational check, not evidence of legal correctness: on 150 newly detected cases, automated records reach 97.2% required-field completeness and cut record-creation time from 218.1 to 1.4 minutes. Table A.7 reports schema completeness and record creation time across three conditions (unstructured baseline, manual CSV entry, and ALERT-automated CSV), corresponding to the RQ4 operational result in the main text.

Table A.7: Schema completeness and record creation time across three conditions. Referenced from the main text (RQ4).

| Condition                  | All Fields (%) | Req. Fields (%) | Time (min) |
|----------------------------|----------------|-----------------|------------|
| Unstructured (baseline)    | 38.2           | 51.4            | 22.3       |
| Manual CSV                 | 74.6           | 91.7            | 18.1       |
| <b>ALERT-automated CSV</b> | <b>81.3</b>    | <b>97.2</b>     | <b>1.4</b> |

Table A.8 reports the component-wise ablation of the agentic loop, showing the contribution of the Goal Checker, Planner, and Retriever query-expansion to risk classification F1 and retrieval Recall@10 (RQ3 in the main text).

Table A.8: Agentic-loop ablation: component-wise contribution to risk classification F1 and retrieval Recall@10 (mean over 5 evaluation runs). Referenced from the main text (RQ3).

| Variant                       | F1          | Recall@10   |
|-------------------------------|-------------|-------------|
| <b>ALERT (Full)</b>           | <b>0.81</b> | <b>0.84</b> |
| – Goal Checker (Single-Agent) | 0.76        | 0.84        |
| – Planner (fixed task order)  | 0.78        | 0.84        |
| – Retriever query-expansion   | 0.79        | 0.76        |

Table A.9 reports Recall@10 on a held-out development set of 200 query–document pairs for three retriever configurations. The hybrid sparse–dense model is adopted as default in all main-paper experiments.

Table A.9: Retriever ablation. Recall@10 on the LegalRisk-RAG development set (200 query–document pairs).

| Retriever Configuration                  | Recall@10   |
|------------------------------------------|-------------|
| Dense (text-embedding-3-large)           | 0.79        |
| LegalBERT (fine-tuned on RECAP)          | 0.81        |
| <b>Hybrid (BM25 + neural re-ranking)</b> | <b>0.84</b> |

## G Reproducibility: Hyperparameters, Prompts, Protocol

### G.1 Goal-Conditioned Agentic Loop (Full Procedure)

Algorithm 2 gives the full ALERT agentic loop summarized in the main text.

#### Algorithm 2 ALERT Goal-Conditioned Agentic Loop

**Require:** Document corpus  $\mathcal{D}$ , product spec  $P$ , goal predicate  $\mathcal{G}$ , evidence threshold  $(k, \tau)$ , max iterations  $T_{\max}$

**Ensure:** Verified risk report  $R^*$ , mitigation guidance  $G^*$

- 1:  $R_0 \leftarrow \emptyset$ ;  $t \leftarrow 0$
- 2: **while**  $\neg \mathcal{G}(R_t) \wedge t < T_{\max}$  **do**
- 3:  $\mathcal{T}_t \leftarrow \text{PLANNER}(R_t, \mathcal{G}, P)$  {Decompose unmet sub-goals}
- 4: **for** each sub-task  $tk \in \mathcal{T}_t$  **do**
- 5:  $E_{tk} \leftarrow \text{RETRIEVER}(tk, \mathcal{D})$  {LegalRisk-RAG}
- 6: **if**  $|E_{tk}| < k \vee \min \text{score}(E_{tk}) < \tau$  **then**
- 7:  $E_{tk} \leftarrow \text{RETRIEVER}(\text{EXPANDQUERY}(tk), \mathcal{D})$
- 8: **end if**
- 9:  $C_{tk} \leftarrow \text{ANALYZER}(E_{tk}, P)$  {Risk score & clustering}
- 10: **end for**
- 11:  $R_{t+1} \leftarrow \text{REPORTER}(\{C_{tk}\}, R_t)$
- 12: **if**  $\text{GOALCHECKER}(R_{t+1}) = \text{TRUE}$  **then**
- 13: **break**
- 14: **end if**
- 15:  $t \leftarrow t + 1$
- 16: **end while**
- 17:  $G^* \leftarrow \text{MITIGATE}(R_t, P)$  {Per-risk action templates}
- 18: **return**  $R^* \leftarrow R_t, G^*$

### G.2 Goal Predicate and Loop Parameters

Table A.10 summarizes the parameters used by Algorithm 2 for all reported experiments.

### G.3 Agent Prompt Templates (Schematic)

Each of the four agents (Planner, Retriever, Analyzer, Reporter) is invoked with a structured prompt of the form:

Table A.10: ALERT goal-predicate and loop hyperparameters used throughout the evaluation.

| Parameter         | Value         | Description                            |
|-------------------|---------------|----------------------------------------|
| $k$               | 3             | Min. retrieved documents per risk item |
| $\tau$            | 0.62          | Min. relevance score (cosine)          |
| $T_{\max}$        | 5             | Max. agentic loop iterations           |
| $\theta$ (LLM)    | GPT-4o, T=0.2 | Controller policy backbone             |
| Embedding dim.    | 3072          | text-embedding-3-large                 |
| Mapping threshold | 0.72          | Bi-encoder product-risk similarity     |

```
[ROLE]: {agent role description}
[INPUTS]: {R_t, P, sub-goal, retrieved evidence}
[GOAL CHECK]: {unmet sub-conditions of G}
[OUTPUT FORMAT]: {structured JSON schema}
```

The full prompt templates, including the few-shot examples of the treatment-edge classifier, are part of the code release (Appendix G, *Code and data release*). The Goal Checker is implemented as deterministic rule evaluation over  $R_t$  (no LLM call), making termination decisions auditable.

#### G.4 Evaluation Protocol Notes

**Dataset split.** The 70/10/20 train/validation/test split is stratified by year and litigation type and held fixed across all baselines (B1–B5 and ALERT (Full)). The 5 random seeds vary only the test-fold sub-sample within the held-out 20% partition.

**Likert evaluation.** Two legal specialists, blind to system identity, rated 75 reports (the same 15 cases for each of the five conditions, so comparisons are paired by case) on a 5-point scale across five dimensions. Reports were presented in randomized order; system identifiers were stripped.

**Statistical tests.** F1 differences in Tables 3 and A.6 are evaluated by a paired bootstrap test (10,000 resamples). Likert differences in RQ2 use the Wilcoxon signed-rank test on per-case mean ratings;  $p$ -values are not corrected for multiple comparisons, and inter-rater agreement was not computed.

**Code and data release.** We will release all code used for the reported results: the sensing pipeline (alert-tracker), treatment-edge extraction and graph construction, the PLRE annotation guidelines and evaluation scripts, and the agentic loop, together with the treatment-weight schedule  $w(\cdot)$ , prompt templates with few-shot examples, split files, seeds, and ALERT-Dataset (CC BY 4.0; code under MIT). For review, an anonymized copy of the code and data is provided as supplementary material; the full public release will follow the re-

Table A.11: Evidence-omission risk under random and temporal calibration splits. Lower is better; values should be read as evidence-coverage risk, not severity-label correctness.

| Split                               | $\alpha = 0.05$ | $\alpha = 0.10$ | $\alpha = 0.20$ |
|-------------------------------------|-----------------|-----------------|-----------------|
| Random resplits                     | 0.043           | 0.082           | 0.171           |
| Temporal 2023–24 $\rightarrow$ 2025 | 0.061           | 0.113           | 0.196           |

Table A.12: Full ALERT system under two backbones (mean  $\pm$  std over runs). Baselines were not re-run on the open-weights stack.

| Configuration                   | Macro F1        | Mean iters    |
|---------------------------------|-----------------|---------------|
| GPT-4o + text-embedding-3-large | $0.81 \pm 0.02$ | $2.4 \pm 0.3$ |
| Llama-3.1-70B + bge-large       | $0.76 \pm 0.02$ | $2.6 \pm 0.3$ |

view period.

**Models, compute, and software.** The controller, Analyzer, and edge classifier use GPT-4o through its API; embeddings use text-embedding-3-large and, for deduplication only, Gemini text-embedding-004. The robustness re-run (Table A.12) uses Llama-3.1-70B (70B parameters) with bge-large, and the retriever ablation uses LegalBERT fine-tuned on RECAP text. We did not log GPU hours or API cost. As an indication of scale, each case requires on average 2.4 loop iterations (at most  $T_{\max}=5$ ), and the test-set evaluations use five runs (seeds as described above). Exact software packages and versions (BM25, re-ranker, bootstrap and conformal-calibration code) are pinned in the dependency file of the released code.

## H Proof of Proposition 1 and Calibration Protocol

Table A.11 reports evidence-omission risk under random and temporal calibration splits: calibrated thresholds meet the budget  $\alpha$  on random splits, while the temporal split exceeds it at  $\alpha \in \{0.05, 0.10\}$ , showing that exchangeability is stressed under temporal shift (RQ6). Table A.12 reports full-system macro F1 and mean loop iterations when the GPT-4o backbone is replaced by an open-weights stack. Only ALERT was re-run on the open-weights stack; baselines were not, so the table does not support cross-backbone comparisons with the baselines.

### H.1 Termination Property (Proof)

Let  $\Phi(R_t) \in \mathbb{Z}_{\geq 0}$  count the unmet atomic sub-conditions of  $\mathcal{G}$  (one per missing severity label,

per under-supported evidence requirement, and per missing mitigation) over the risk items in  $R_t$ . By assumption (A1) *monotone retrieval*, the evidence set attached to any item is non-decreasing across iterations, so a sub-condition once satisfied by  $\mathcal{G}_{\text{evidence}}$  or  $\mathcal{G}_{\text{guide}}$  remains satisfied; hence no satisfied atom can become unmet and  $\Phi(R_{t+1}) \leq \Phi(R_t)$ . By assumption (A2) *attainable progress*, whenever  $\Phi(R_t) > 0$  and at least one unmet atom is *attainable* (the required evidence exists in  $\mathcal{D}$ ), the Planner schedules it and the Retriever/Analyzer resolves it, giving  $\Phi(R_{t+1}) \leq \Phi(R_t) - 1$ . Therefore, if every atom is attainable,  $\Phi$  strictly decreases each iteration until  $\Phi = 0$ , i.e.  $\mathcal{G} = \text{TRUE}$ , in at most  $\Phi(R_0)$  steps. If some atom is unattainable,  $\Phi$  cannot reach 0; the loop runs to  $T_{\max}$  and the Goal Checker reports the residual  $\Phi(R_{T_{\max}})$  unmet atoms, which are exactly the items requiring human escalation.  $\square$

*Remark.* (A1) holds by construction because EXPANDQUERY unions new hits with the prior evidence cache; (A2) is the design assumption that one tractable sub-goal is closed per pass. We do not verify it directly; the mean iteration count (2.4, below  $T_{\max}=5$ ) is consistent with it but does not establish it.

## H.2 Conformal Risk-Controlled Termination

We calibrate the evidence threshold  $\tau$  (and, jointly,  $k$ ) so that termination meets a user-specified omission-risk budget  $\alpha$ . For calibration cases  $c_i$ ,  $i = 1, \dots, n$ , let the bounded loss  $L_i(\tau)$  indicate that a retrievable, human-labeled High-risk theory of  $c_i$  is omitted from  $R^*(c_i)$ , and let  $\hat{R}_n(\tau) = \frac{1}{n} \sum_i L_i(\tau)$ , which is non-increasing as  $\tau$  decreases (more evidence admitted). Following conformal risk control (Angelopoulos et al., 2023), choosing  $\hat{\tau} = \inf\{\tau : \frac{n}{n+1} \hat{R}_n(\tau) + \frac{1}{n+1} \leq \alpha\}$  controls  $\mathbb{E}[L_{n+1}(\hat{\tau})] \leq \alpha$  for an exchangeable test case under the calibration assumptions. For the joint  $(k, \tau)$  we use a Learn-then-Test grid (Angelopoulos et al., 2021), treating each configuration as a hypothesis and controlling the family-wise error via fixed-sequence testing. The achieved-vs-target risk is reported in RQ6 and Table A.11.

## H.3 Temporal-Split Evidence-Coverage Audit

Table A.11 separates the standard random calibration/test split from a stricter temporal split. The temporal split does not invalidate the calibrated procedure; it shows that the exchangeability assumption is empirically stressed by evolving liti-

gation streams and motivates rolling or weighted calibration in deployment.

## I Pre-Specified Validation Protocols

The two construct-validity steps flagged in the Discussion—transfer of the validity-gating modifier to an external retrieval stack, and agreement of the automatically extracted treatment graph with an authoritative citator—are specified here as *pre-specified protocols*. Protocol P1 (external-stack transfer) is reported as a protocol only and is not yet executed; Protocol P2 (citator agreement) has been executed and its measured results are reported in Appendix I.2. The protocols fix the datasets, conditions, metrics, statistical tests, and acceptance criteria *before* execution so that the outcome cannot be tuned post hoc. Both protocols are designed to be runnable on public resources plus a single licensed citator account, and both release their analysis scripts, query/seed manifests, and (for the citator study) the category crosswalk rather than any redistributed proprietary content.

### I.1 Protocol P1: External-Stack Transfer of the Validity-Gating Modifier

**Hypothesis.** The supersession-aware score of Eq. (3)—a base similarity multiplied by a time-validity indicator  $\mathbf{1}[s_v \leq t_q < e_v]$  and a treatment weight  $w(\text{treat}(\cdot))$ —improves *external* legal retrievers that do not use ALERT’s agentic loop, dataset, or goal predicate, which would indicate that the mechanism transfers rather than being ALERT-specific.

**Host retrievers (unmodified, external).** (i) BM25; (ii) a general dense retriever (text-embedding-3-large or an open bge/e5 encoder); (iii) a legal encoder (LegalBERT). Each is the “host” stack; the modifier is applied only as a post-hoc re-scoring layer ( $\text{score} = \text{score}_{\text{host}} \cdot \mathbf{1}[\cdot] \cdot w(\cdot)$ ), with *no* re-indexing and *no* ALERT loop.

**Evaluation corpora (disjoint from ALERT-Dataset).** (a) a COLIEE-style public case-retrieval benchmark; and (b) a temporally held-out CourtListener slice whose jurisdiction/time window does not overlap the ALERT-Dataset window. Setting (b) is the decisive transfer test because it contains time-varying authority outside ALERT’s training window.

**Conditions.** For each host retriever: {ungated baseline} vs. {+ indicator only} vs. {+ indicator

and treatment weight}, mirroring the decomposition of Table 4. Each query carries a query date  $t_q$ ; an authority closed before  $t_q$  is treated as stale (non-relevant) in the gold judgments.

**Treatment-label source (confound control).** To avoid inheriting ALERT’s extractor noise, the external test uses treatment labels from a source independent of ALERT’s extractor (gold/citator-derived). We report two variants per condition—*gold-treatment* and *extractor-treatment*—so the transfer claim is separated from extraction quality.

**Metrics.** Retrieval quality (nDCG@10, Recall@10, MRR) and a *stale-citation rate* (fraction of top- $k$  closed before  $t_q$ ). The primary endpoints are the paired deltas (gated – ungated) per host retriever.

**Statistics.** Paired bootstrap over queries (10,000 resamples) with 95% CIs on each delta; Holm correction across the three host retrievers; a pre-specified non-inferiority margin ( $\delta_{NI} = 0.5$  nDCG points) requiring the modifier not to hurt on queries with no negative-treatment in scope.

**Acceptance criterion (pre-specified).** P1 succeeds iff the modifier yields a positive nDCG@10 delta *and* a reduced stale-citation rate on at least two of three host retrievers, the effect persists on the temporally held-out slice (b), and non-inferiority holds on the no-stale subset. Failure on slice (b) but success on (a) is reported as “in-distribution only,” which would *not* support the transfer claim.

## I.2 Protocol P2: Citator Agreement for the Treatment Graph

**Objective.** Replace “validated against our own gold spans” with “validated against an authoritative citator” for the negative-treatment edges (SUPERSEDES, OVERRULES, material NARROWS) that drive interval closure.

**Sampling frame and power.** A stratified random sample of edges from the current graph (1,812 nodes, 4,760 edges), stratified by relation type and polarity and oversampling negative-treatment edges. Target  $n = 400$  edges with  $\geq 150$  negative-treatment edges; this supports a 95% CI half-width of  $\leq 0.07$  on Cohen’s  $\kappa$  at an anticipated  $\kappa \approx 0.75$ . The exact draw is fixed by a released seed before any citator lookup.

**Reference standard.** For each sampled (citing  $\rightarrow$  cited) pair and date, retrieve the citator treatment signal (e.g., KeyCite/Shepard’s negative, cau-

tionary, or positive flags) and map it to ALERT’s relation/polarity taxonomy through a *pre-specified, published crosswalk*. Two annotators perform this mapping independently and blind to ALERT’s predicted edge; a third adjudicates disagreements. Inter-annotator agreement on the citator mapping itself is reported (Cohen’s  $\kappa$ ).

### Primary endpoints.

- *Edge agreement:* precision/recall/F1 of ALERT edges vs. citator-derived edges, per relation and micro-averaged over negative treatment.
- *Interval-closure agreement:* whether ALERT closes an interval no later than the citator’s negative-treatment date; the closure-date error distribution (in days); and the rates of *missed closures* (citator negative, ALERT open) and *false closures* (ALERT closed, citator still good).
- *Point-in-time correctness:* validity intervals are checked at multiple historical dates, not only the latest status, because citators are themselves time-varying.

**Downstream link.** On the disagreement subset, the PLRE/stale-cite audit is re-run with citator-corrected closures to measure the *decision-flip rate*, connecting construct validity to risk-decision outcomes.

**Acceptance criterion (pre-specified).** P2 succeeds iff negative-treatment micro-F1 vs. the citator  $\geq 0.80$  and the missed-closure rate is at or below the extractor’s audited miss rate (23%) underlying Table 5; relation types failing the bar are routed to the human review queue and the residual is reported rather than suppressed.

**Licensing and ethics.** Citator outputs are proprietary and are used solely as an evaluation reference under a licensed account. The release contains the crosswalk, sampling seeds, and agreement statistics, but no redistributed citator content, consistent with the data-governance boundary in the main text.

**Measured P2 results.** The audit sampled  $n = 400$  treatment candidates, including 162 negative-treatment candidates. Inter-annotator agreement on the citator-to-ALERT crosswalk was substantial ( $\kappa = 0.78$ ). Edge-level agreement was measured against the citator-derived reference, and interval-closure behavior plus downstream impact were reported. Negative-treatment micro-F1 is 0.81 and the missed-closure rate is 20.4%, satisfying the pre-specified P2 threshold and staying below the

Table A.13: Measured P2 edge agreement: ALERT treatment edges vs. the authoritative citator-derived reference. Micro-average is over the three negative-treatment relations.

| Relation                          | <i>n</i> | Precision   | Recall      | F1          |
|-----------------------------------|----------|-------------|-------------|-------------|
| OVERRULES                         | 54       | 0.86        | 0.82        | 0.84        |
| SUPERSEDES                        | 47       | 0.82        | 0.79        | 0.80        |
| material NARROWS                  | 61       | 0.79        | 0.76        | 0.77        |
| <b>Negative-treatment (micro)</b> | 162      | <b>0.83</b> | <b>0.80</b> | <b>0.81</b> |
| Cautionary treatment              | 78       | 0.81        | 0.74        | 0.77        |
| Positive/neutral citation         | 160      | 0.93        | 0.95        | 0.94        |

Table A.14: Measured P2 closure agreement and downstream decision impact against the citator-derived reference.

| Metric                                           | Value     |
|--------------------------------------------------|-----------|
| Matched negative closures                        | 129 / 162 |
| Missed-closure rate                              | 20.4%     |
| False-closure rate                               | 4.8%      |
| Median closure-date error                        | 18 days   |
| Mean closure-date error                          | 31 days   |
| 90th-percentile closure-date error               | 74 days   |
| Point-in-time validity accuracy                  | 0.88      |
| PLRE decision-flip rate after citator correction | 3.8%      |

$\approx 23\%$  internal-gold miss rate of the extractor configuration evaluated in Table 5, so the measured downstream effect there is not optimistic relative to the citator reference. Material NARROWS remains the weakest relation and is retained in the human-review queue. Edge-level agreement is reported in Table A.13; interval-closure agreement and downstream decision impact are reported in Table A.14.

1495  
 1496  
 1497  
 1498  
 1499  
 1500  
 1501  
 1502  
 1503